

(18)

MAPREDUCE TYPES AND FORMATS

Hadoop uses the MapReduce programming model for the data processing of input and output for the map and to reduce functions represented as key-value pairs. They are subject to parallel execution of datasets situated in a wide array of machines in a distributed architecture. The programming paradigm is essentially functional in nature in combining while using the technique of map and reduce. This article introduces the MapReduce model, and in particular, how data in various formats, from simple text to structured-binary objects are used.

MAPREDUCE TYPES.

Mapping is the core technique of processing a list of data elements that come in pairs of keys and values. The map functions applies to individual elements defined as key-value-pairs of a list and produces a new list. The general idea of map and reduce function of Hadoop can be illustrated as follows:-

map: $(K_1, V_1) \rightarrow \text{list}(K_2, V_2)$

reduce: $(K_2, \text{list}(V_2)) \rightarrow \text{list}(K_3, V_3)$

The Default MapReduce Job.

The default mapper is IdentityMapper, which writes the input key and value unchanged to the output.

- Mapreduce take Two arguments (input & output)

- Public class IdentityMapper $\langle K, V \rangle$
extends MapReduceBase implements Mapper $\langle K,$
 $V, K, V \rangle$

```
{  
    public void map (K key, V val, output collector  
                     $\langle K, V \rangle$  output, Reporter, reporter)
```

```
    throws IOException
```

```
{
```

```
    output.collect(key, value;
```

```
}
```

```
}
```

- IdentityMapper is a generic type, which allows it to work with any key or value types, with the restriction that the map input and output keys are same type, and map input and output values are same type.

- map output key is LongWritable.

- map output value is Text.

* Input Formats

Hadoop can process many different types of data formats, from flat text files to data bases.

- Different formats are available.

* Input Splits and Records

Input split is a chunk of the input that is processed by a single map. The `InputSplit` represents the data to be processed by a map. It returns the length in bytes and has a reference to the input data. It presents a byte-oriented view on the input and is the responsibility of the `RecordReader` of the job to process this and present a record-oriented view.

- In most cases, we do not deal with `InputSplit` directly because they are created by an `InputFormat`.

eg: Public Interface `InputFormat <K, V> {`

`InputSplit [] getSplits (Job conf job,`

`IO Exception;`

`RecordReader <K, V> getRecordReader (Input-`

`Job conf job, throws IO Exception;`

`}`

- * The `FileInputFormat` is the base class for the file-data source. It has the responsibility to identify

the files that are to be included as the job input and the definition for generating the split.

- * Hadoop also includes processing of unstructured data that often comes in internal formats. The TextInputFormat is the default input format for such data.
- * The SequenceInputFormat takes up binary inputs and stores sequences of binary key-value pairs.
- * DatabaseInputFormat provides the capability to read data from relational database using JDBC.

Text Input

Hadoop excels at processing unstructured text.

TextInputFormat

It is the default input format. Each record is a line of input. The key, a longWritable is the byte offset within the file of the beginning of the line. The value is the contents of the line, excluding any line terminators and is packaged as a text object.

KeyValue TextInputFormat

It's keys being simply the offset within the file are not normally very useful. It is -

common for each line in a file to be a key-value pair, separated by a delimiter such as a tab character.

LineInputFormat

With TextInputFormat and KeyValueTextInputFormat, each mapper receives a variable no. of lines of input. The number depends on the size of the split and the length of the lines.

XML

- Most XML parsers operate on whole XML documents so if a large XML document is made up of multiple input splits - then it is a challenge to parse these individually.
- Hadoop comes with a class for this purpose - called StreamXMLRecordReader.

Binary Input

Hadoop MapReduce is not just restricted to processing text data - it has support for binary formats too.

SequenceFile Input Format

Hadoop's Sequence File format stores sequences of binary key-value pairs. They are well suited as a format for MapReduce data.

since they are splittable, they support chaining as a part of the format and they can store arbitrary type using variety of serialization Framework.

Multiple Inputs

Although the input to a MapReduce job may consist of multiple input files all of the input is interpreted by single InputFormat and a single Mapper.

Database Input (& output)

DBInputFormat is an InputFormat for reading data from a relational database using JDBC. Because it doesn't have any sharding capabilities you need to be careful not to overwhelm the database you are reading from by running too many mappers.

SequenceFile as Input Format

It is a variant of SequenceFileInputFormat that converts the sequence file's key and values to Text Objects. The conversion is performed by calling to string() on the keys and values. This format makes sequence file suitable input for streaming.

OUTPUT FORMATS

The output format classes are similar to their corresponding input format classes and work in the reverse direction.

* TextOutputFormats

It is the default output format that writes records as plaintext files, whereas key values can be any ~~of~~ types, and transforms them into a string by working ~~they~~ to strings method. The key-value character is separated by the tab character, although this can be customized by manipulating the separator property of the text output formats.

* BinaryOutput

There is sequence file input format to write a sequence of binary output to a file. Binary outputs are particularly useful if the output becomes input to a further mapReduce job.

* The output formats for relational databases and to HBase are handled by DB output format.

* Multiple Outputs

It allows writing data to files whose names are derived from output keys and values or in fact from an arbitrary string.

Lazy output formats

Sometimes file output formats will create output files even if they are empty. Lazy-output formats is wrapper output formats which ensures that the output file will be created only when the record is emitted for given position.

* MAPREDUCE FEATURES..

① Counters

Hadoop counters provides a way to measure the progress or the no. of operations that occur within map/reduce job. Counters in Hadoop-MapReduce are a useful channel for gathering statistics about the mapreduce job, for quality control or for application-level. They are also useful for problem diagnosis.

- Counters represents Hadoop Global counters, defined either by the mapreduce framework or applications.
- Counters are branched into groups. each comprising of counters from a particular Enum class.

* Hadoop candidates validate that :-

- The correct no. of bytes was read and written.
- The correct no. of tasks was launched & successfully

The amount of CPU and memory consumed is appropriate for our jobs and cluster nodes.

2) Built-In counters

Hadoop maintains some built-in hadoop counters for every job and these report various metrics, like, there are counters for the number of bytes and records, which allows us to confirm that the expected amount of input is consumed and the expected amount of output is produced.

- Hadoop counters are divided into groups and there are several groups for the built-in counters. Each group either contains task counters or job counter.
- There are several groups for the hadoop built-in counters.
- Hadoop task counters are maintained by each task attempt and periodically sent to the application master so they can be globally aggregated.

3) User-Defined Java Counters

Hadoop MapReduce permits user code to define a set of counters. Then it increments them as desired in the mapper or reducer. Like in Java to define counters is user, 'enum'.

- A job may define an arbitrary number of 'enums' each with an arbitrary number of fields. The

*
name of the enum is the group name. The enum's fields are the counter names.

i) Dynamic counters in Hadoop

Java enum's fields are defined at compile time. So we cannot create new counters at runtime using enums. So, we use dynamic counters to create new counters at runtime. But dynamic counter is not defined at compile time.

ii) Readable counter names

By default, a counter's name is the enum, fully qualified Java classname. These names are not very readable when they appear on the web UI or in the console, so Hadoop provides a way to change the display names using resource bundles.

iii) Retrieving counters

In addition, to being available via the web UI and the command line, you can retrieve counter values using the Java API.

iv) User-defined Streaming Counters

A streaming MapReduce program can increment counters by sending a specially formatted line to the standard error stream, which is co-opted as a control channel in this case.

* Sorting

The ability to sort data is at the heart of MapReduce. Even if your application isn't concerned with sorting per se, it may be able to use the sorting stage that MapReduce provides to organise its data.

i) Preparation

We are going to sort the weather dataset by temperature. Storing temperatures as Text-objects doesn't work for sorting purposes, since signed integers don't sort lexicographically.

ii) Partial Sort

The reducer output will be lot of files each of which is sorted within itself based on the key.

- N number of mappers will simply generate N number of reducers will sort these files individually.

iii) Total Sort

- The reducer output will be a single file having all the output sorted based on the key.
- All key value pairs from a particular key will reach a particular reducer. This will happen through partitioners at mapper level. Combiners at mapper level will act as semi-reducers and

send values of a particular key to reducer.

iv) Secondary sort

We will be able to control the ordering of the values along with the keys. That is sorting can be done on two or more field values.

* Joins

MapReduce can perform join between large-dataset, but writing the code to do join from-scratch is fairly involved. Rather than writing-MapReduce prgm, you might consider using a higher-level framework such as pig, Hive or - cascading in which join operation are a core part of the implementation.

- We have two datasets.
for eg: the weather station database and the weather records - we want to reconcile the two.
- How we implement the join depends on how large the datasets and how they are partitioned.
- If one dataset is large but the other one is - small enough to be distributed to each node in the cluster.

* Side data distribution

Side data distribut can be defined as extra read only data needed by a job to process the main dataset. The challenge is to make side-data available to all the map or reduce-tasks in a convenient and efficient function.

* Distributed cache

Rather than serializing side data in the job configuration, it is preferable to distribute dataset using Hadoop's distributed cache mechanism. This provides a service for copying files and archives to the task ~~map~~ nodes in-time for the task to use them when they run. To save network bandwidths, files are normally copied to any particular node once per job.

At
04/11/201

